



# PyTorch and CNNs



# Agenda

1. Overview of AI and Deep Learning (17.15-17.25)
2. Datasets (17.25-17.30)
  - a. Task: Implement Dataset in Pytorch (17.30-17.40)
  - b. Solution (17.40-17.45)
3. CNNs
  - a. Convolution (17.45-18.00). Task (18.15-18.25). Solution (18.25-18.30)
  - b. Other Layers (18.30-18.35). Task (18.35 - 1845). Solution (1845-1850) (DROP)
4. Training Loop (18.50 - 19.00). Task (19.00-19.15). Solution (19.15-1920)
5. Specific CNN Architectures (19.20-19.25).
6. Competition: Create the most accuracy custom CNN! (19.25-20.00)

# HELLO!

I AM ULRIK RØSBY

- 5. Year Computer Science
- Writing master's in Autonomous Driving
- Earlier been Project lead in 2 projects at Cogito NTNU (Emojify, CatMatch)





1

# Introduction

Welcome to deep learning



# What is AI?

- Defining AI is hard
  - John McCarthy:
    - “It is the science and engineering of making intelligent machines, especially intelligent computer programs.”

What is intelligence?

- Doing the correct thing or something like that

Quote source: <http://jmc.stanford.edu/artificial-intelligence/what-is-ai/index.html>

Image source:

[https://upload.wikimedia.org/wikipedia/commons/thumb/9/9d/Artificial\\_intelligence\\_Reasoning%2C\\_problem-solving.jpg/800px-Artificial\\_intelligence\\_Reasoning%2C\\_problem-solving.jpg](https://upload.wikimedia.org/wikipedia/commons/thumb/9/9d/Artificial_intelligence_Reasoning%2C_problem-solving.jpg/800px-Artificial_intelligence_Reasoning%2C_problem-solving.jpg)





# What is Deep Learning?

If statement

Linear regression

Neural networks

Artificial Intelligence

Machine Learning

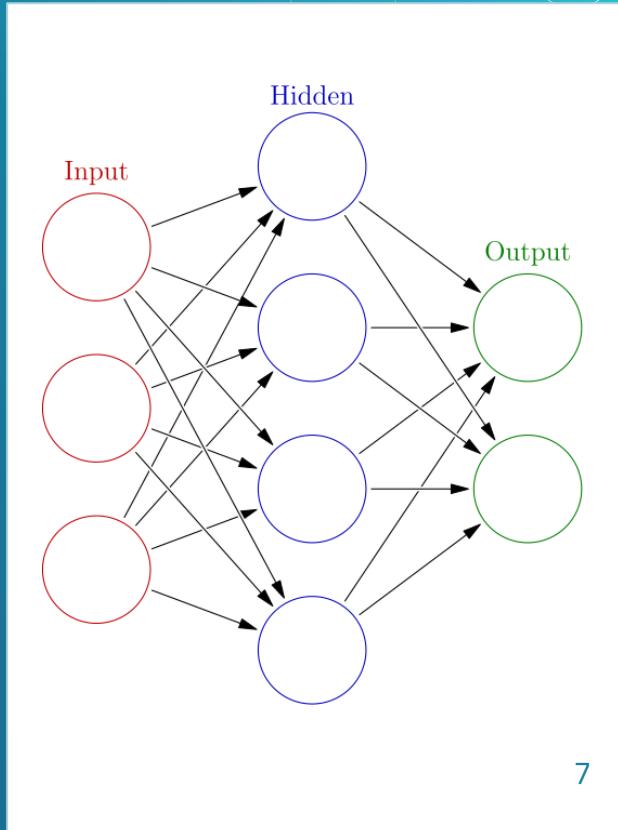
Deep Learning



# What is a Neural Network?

- Often use Artificial Neural Network (ANN)
- Consists of 3 parts
  - Input layer
  - Hidden layers
  - Output layer
- Each layer consists of artificial neurons
  - $a(x) = f(w \cdot x + b)$
  - $x$  - input
  - $w$  - learned weights
  - $b$  - learned bias
  - $f$  - activation function (chosen beforehand)

Image source: [https://upload.wikimedia.org/wikipedia/commons/thumb/4/46/Colored\\_neural\\_network.svg/1200px-Colored\\_neural\\_network.svg.png](https://upload.wikimedia.org/wikipedia/commons/thumb/4/46/Colored_neural_network.svg/1200px-Colored_neural_network.svg.png)





# Training a neural network

- 3 main methods
  - **Supervised Learning**
    - We know the correct output
  - **Unsupervised learning**
    - We have no correct output
  - **Reinforcement learning**
    - Give the model a reward based on how correct it is
- *We will only care about **supervised learning** today*

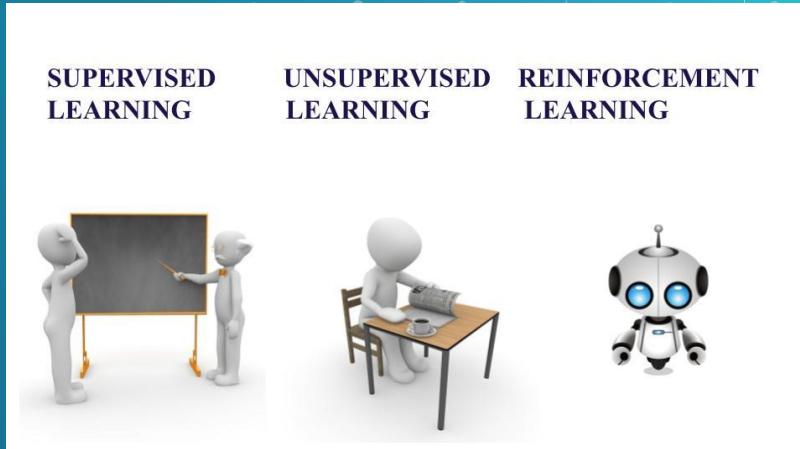


Image source: <https://www.aitude.com/wp-content/uploads/2020/01/Untitled-presentation-1-1.jpg>



# What is a Neural Network? (2)

Steps for training a ANN with supervised learning

1. Randomly initialize the weights  $\mathbf{w}$  and bias  $\mathbf{b}$ .
2. Do a *forward pass* through the network  
(calculate the result for an input  $\mathbf{x}$ ). Thus obtain the predicted result  $\hat{\mathbf{y}}$ .
3. Using the loss function  $\mathbf{L}$ , calculate  $\mathbf{L}(\hat{\mathbf{y}}, \mathbf{y})$
4. Do backpropagation to find the gradient in the network
5. Use the optimizer, like SGD, to use the gradient to optimize the weights
6. Repeat from 2. until satisfied

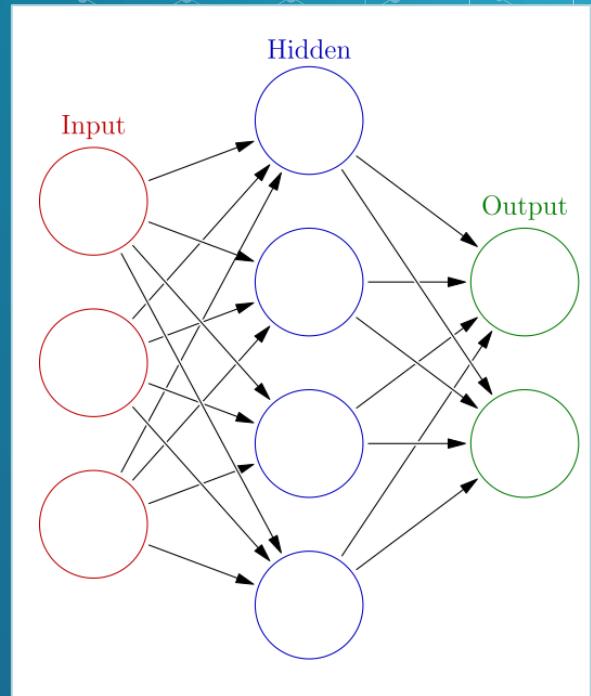


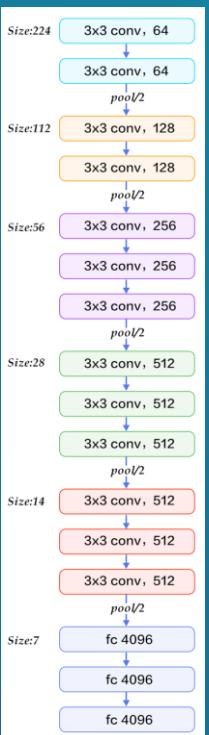
Image source: [https://upload.wikimedia.org/wikipedia/commons/thumb/4/46/Colored\\_neural\\_network.svg/1200px-Colored\\_neural\\_network.svg.png](https://upload.wikimedia.org/wikipedia/commons/thumb/4/46/Colored_neural_network.svg/1200px-Colored_neural_network.svg.png)



# Architecture of a Neural Network

## Architecture of VGG16 a CNN

Image source: <https://www.kaggle.com/code/blurredmachine/vggnet-16-architecture-a-complete-guide>





# What is a CNN?

- It is a method to analyze images in a neural network
- Groups parts of the image to produce a smaller representation called a **feature map**

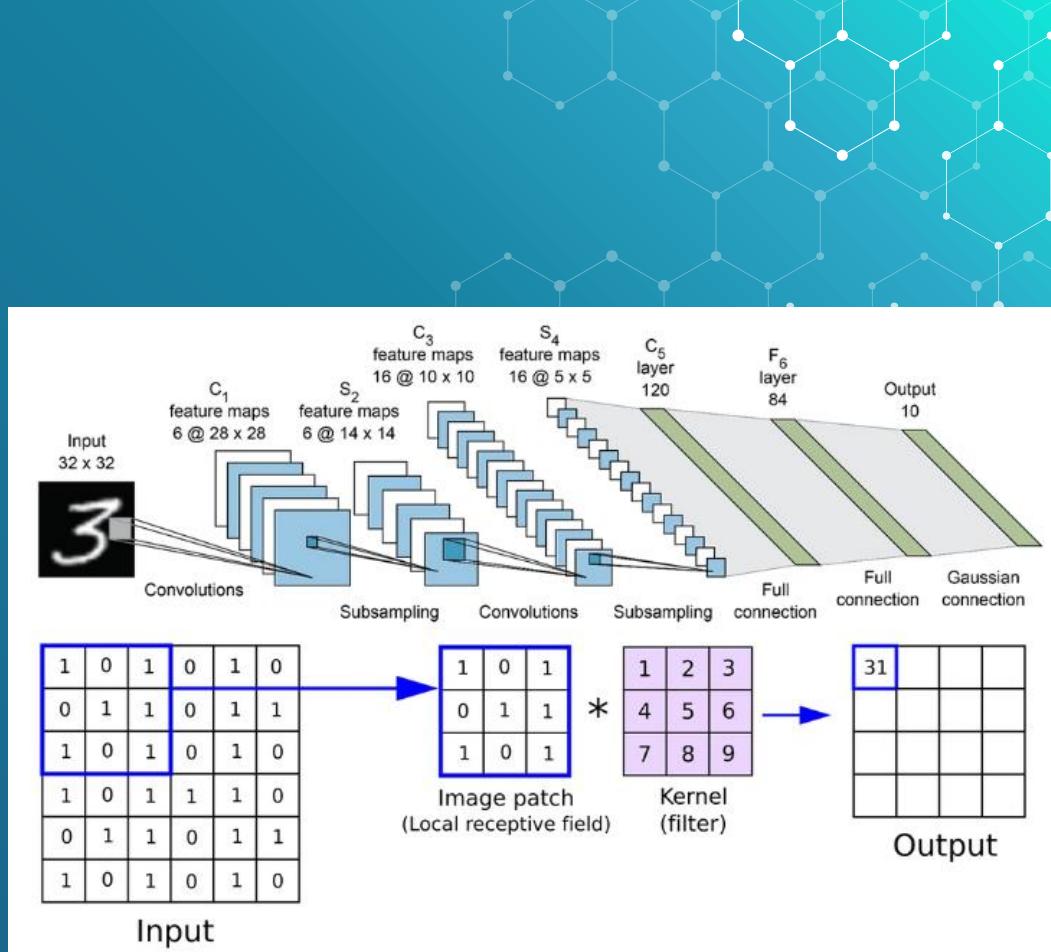
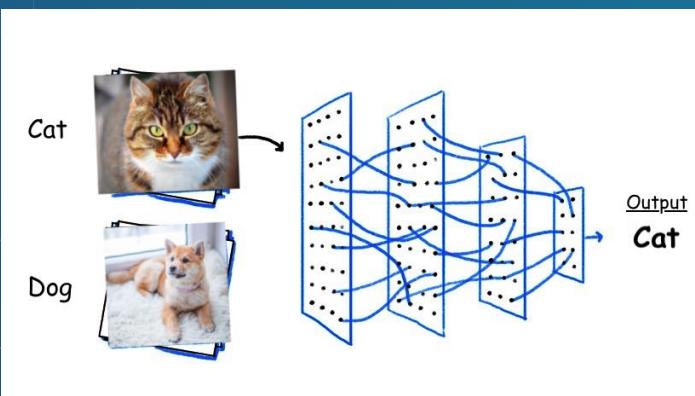


Image source: [https://assets-global.website-files.com/614c82ed388d53640613982e/646371e3bdc5ca90dee5331b\\_convolutional-neural-network%20\(1\).webp](https://assets-global.website-files.com/614c82ed388d53640613982e/646371e3bdc5ca90dee5331b_convolutional-neural-network%20(1).webp)

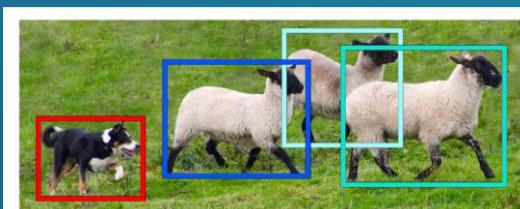


# What can a CNN do? - examples

## Classification



## Object Detection



## Semantic Segmentation

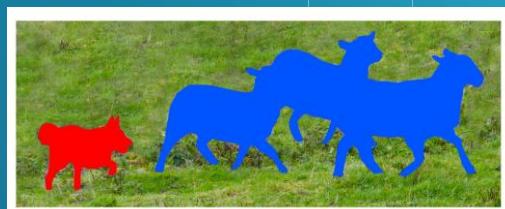


Image sources:

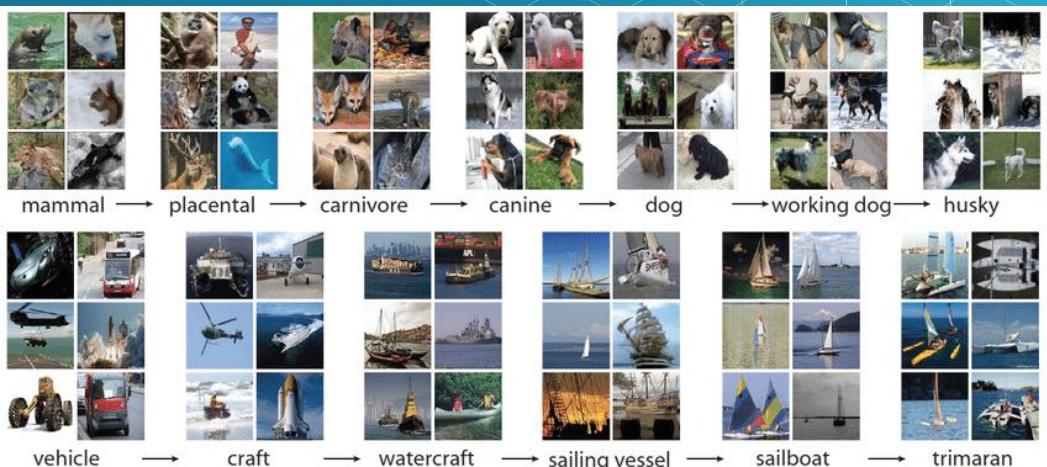
Left: <https://towardsdatascience.com/10-papers-you-should-read-to-understand-image-classification-in-the-deep-learning-era-4b9d792f4a7>  
Middle and right: <https://manipulation.csail.mit.edu/segmentation.html>



# Datasets

- Is a set of data that is used to train the model
- Especially used in supervised learning to ensure having the “correct” values - *ground truth*

ImageNet (14 mill labeled images).





# Typical datasets

- Premade datasets  
(ImageNet, CIFAR100, Cityscapes, nuScenes...)
- CSV files
- Find open datasets on <https://kaggle.com>



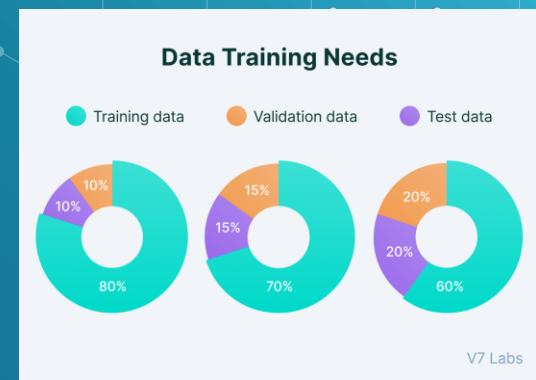
# Important aspects of datasets

- Balanced distribution of data
- Include all relevant situations
- Large (usually several 1000 elements)
- No annotation errors



# Processing of dataset during training

- The dataset is split into
  - Training dataset (often 80%)
  - (Validation dataset) (often 10%)
  - Testing dataset (10% - 20%)
- Processed in batches
- Format the data to a way the model efficiently understands





# PyTorch

- Is research industry standard library for deep learning
  - Includes all you tools you need to build, train and run a deep learning model
- Also includes some premade and pretrained models in torchvision
- Alternatives: Tensorflow, JAX



Image source:  
[https://github.com/pytorch/pytorch/blob/main/docs/source/\\_static/img/pytorch-logo-dark.png](https://github.com/pytorch/pytorch/blob/main/docs/source/_static/img/pytorch-logo-dark.png)



2

# Datasets

What the model is learning from



# HOW TO REPRESENT DATA?



# How to represent data?

- Images
  - Each pixel is a vector of their red, green and blue value [255, 165, 0]
  - Thus a image that is 512 \* 256 px has shape (512, 256, 3)
  
- Labels
  - Numbers
  - If your labels are ["cat", "dog", "mouse"]
    - The labels are cat=1, dog=2, mouse=3 and are represented by position in an array in the network



# Datasets in PyTorch

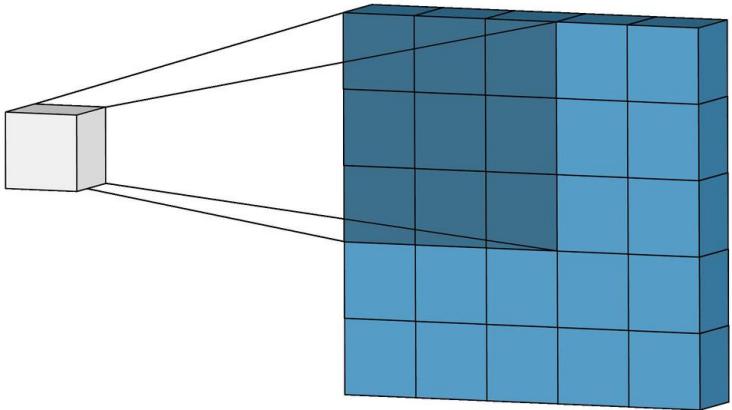
- On to the Google Colab Notebook



3

# CNNs

## Convolutional Neural Networks



# CONVOLUTION



# Convolution - the continuous math def

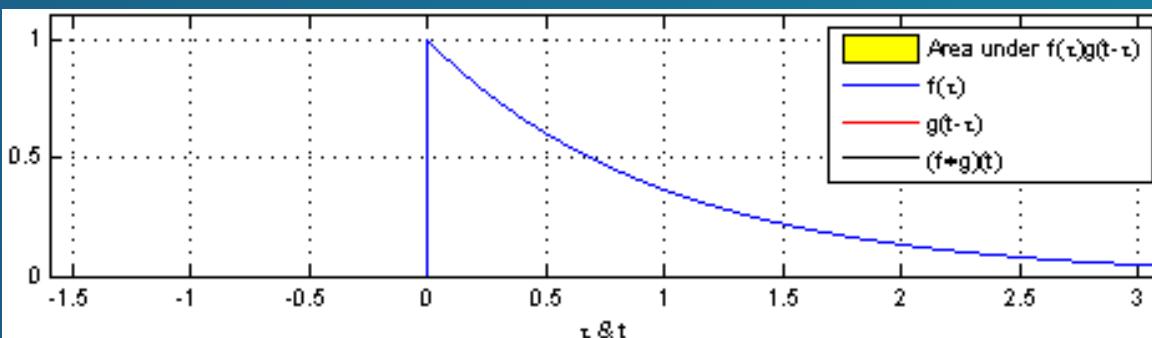
$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$



# Convolution - mathematical definition

A way to combine two functions

$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$

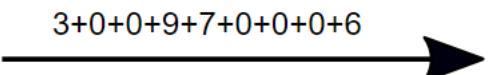




# Convolution - on images

Input Feature Map

3x1	5x0	2x0	8	1
9x1	7x1	5x0	4	3
2x0	0x0	6x1	1	6
6	3	7	9	2
1	4	9	5	1



Output Feature Map

25	18	17
18	22	14
20	15	23

Convolutional Filter

1	0	0
1	1	0
0	0	1

Image source: <https://developers.google.com/machine-learning/practica/image-classification/convolutional-neural-networks>



# Convolution - on images

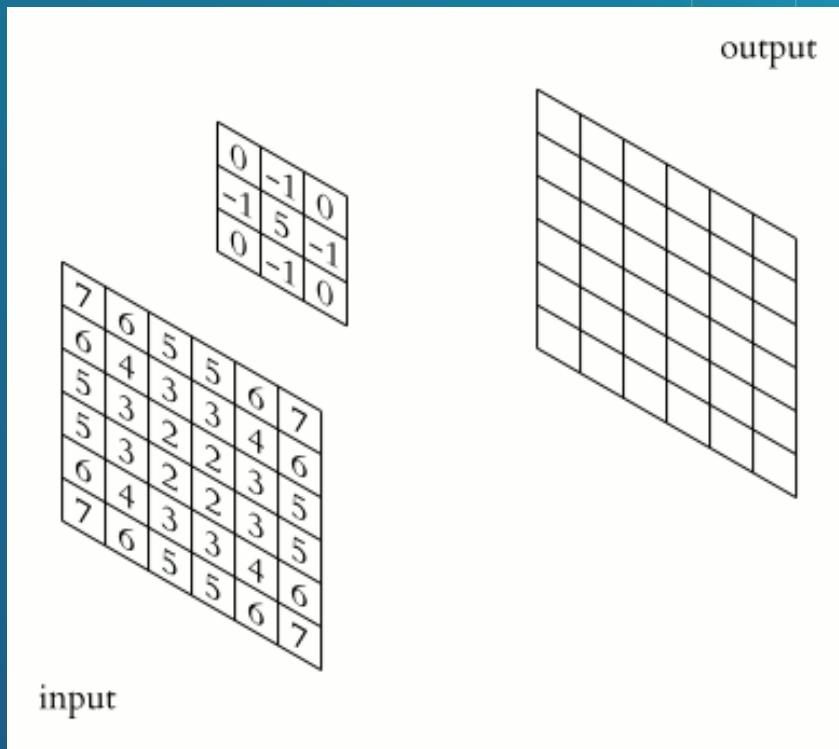


Image source:  
[https://upload.wikimedia.org/wikipedia/commons/thumb/1/19/2D\\_Convolution\\_Animation.gif/330px-2D\\_Convolution\\_Animation.gif](https://upload.wikimedia.org/wikipedia/commons/thumb/1/19/2D_Convolution_Animation.gif/330px-2D_Convolution_Animation.gif)



# Exercise!

- Implement a convolution on an image!



# Convolutional Layers

- We are trying to find the kernel!
- The weights in a convolutional layer denote the numbers that make up the kernel



# Convolutional Neural Networks

- A CNN is an ANN with at least one convolutional layer
- Other layers:
  - Batch Normalization (BatchNorm)
  - Max Pool
- A CNN outputs a feature map
  - A lower-dimensionality representation of the image



# Max Pooling

- Uses a  $n \times n$  kernel and a stride
- Goes through the matrix like convolution,
- But chooses the maximum value instead

2	2	7	3
9	4	6	1
8	5	2	4
3	1	2	6

Max Pool  
→

Filter - (2 x 2)  
Stride - (2, 2)

9	7
8	6



# Max Pooling - Task!

- Let's implement Max Pooling!



# Batch Normalization

- Uses re-centering and scaling to make ANNs faster and more stable
- It normalizes the distribution of output values from the neurons
- This way large values don't impact the result that much

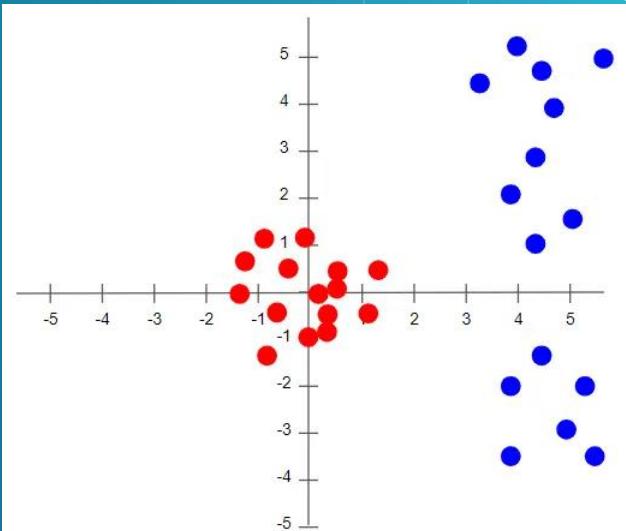


Image source: <https://towardsdatascience.com/batch-norm-explained-visually-how-it-works-and-why-neural-networks-need-it-b1891969279>



4

# Training Loop

The most successful ones

# Neural Networks - Activation Functions

- Introduces non-linearity into the network

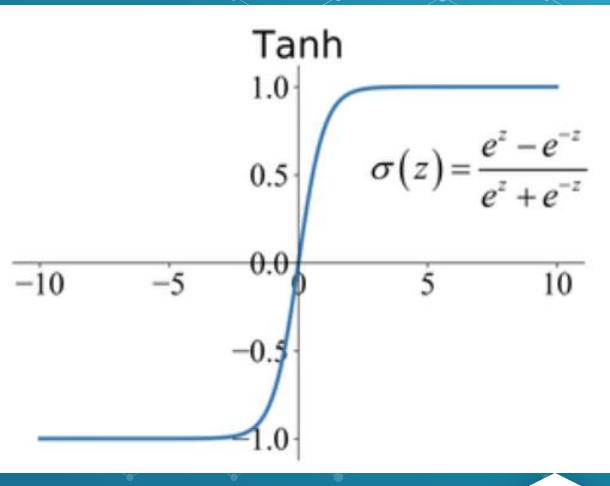
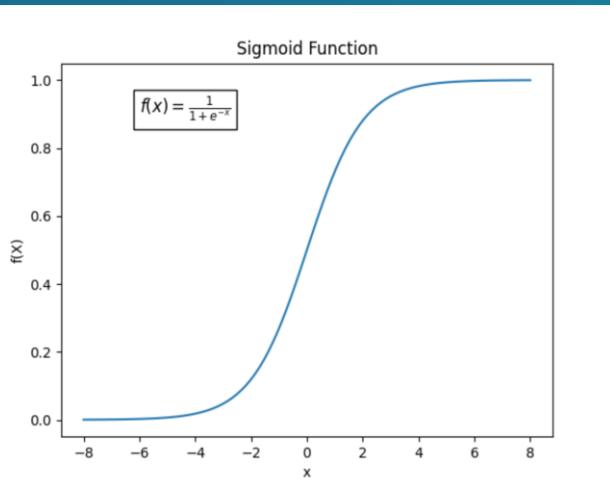
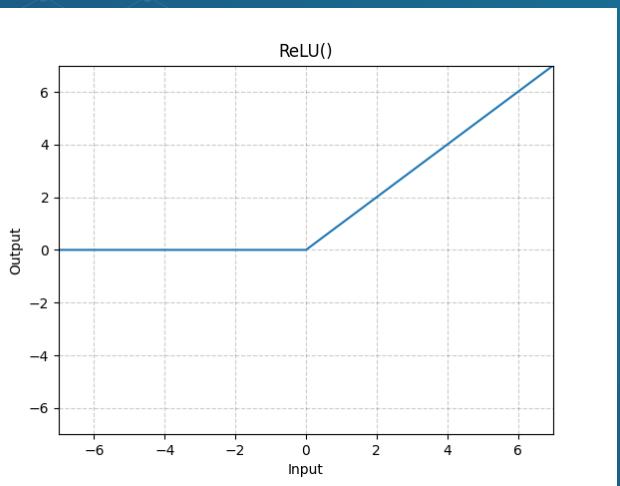


Image source:  
[Left: https://pytorch.org/docs/stable/\\_images/RelU.png](https://pytorch.org/docs/stable/_images/RelU.png)

Middle: <https://raw.githubusercontent.com/Codecademy/docs/main/media/sigmoid-function.png>

Right: [https://production-media.paperswithcode.com/methods/Screen\\_Shot\\_2020-05-27\\_at\\_4.23.22\\_PM\\_dcuMBJl.png](https://production-media.paperswithcode.com/methods/Screen_Shot_2020-05-27_at_4.23.22_PM_dcuMBJl.png)

# Neural Networks - Loss

- Are function that calculate how “wrong” the prediction is.
- Their derivative decide the gradient
- Examples (for regression):

## Mean Squared Error

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

Image sources:

Left: en.wikipedia.org/wiki/Mean\_squared\_error

Right: <https://editor.analyticsvidhya.com/uploads/42439Screenshot%202021-10-26%20at%209.34.08%20PM.png>

Steps for training a ANN with supervised learning

1. Randomly initialize  $\mathbf{w}$  and  $\mathbf{b}$ .
2. Do a *forward pass* to obtain  $\hat{\mathbf{y}}$ .
3. Calculate loss:  $L(\hat{\mathbf{y}}, \mathbf{y})$
4. Do backpropagation
5. Use the optimizer (e.g.) SGD to improve weights.
6. Repeat from 2. until satisfied

## Mean Absolute Error

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

# Neural Networks - Loss (2)

- For classification

## Cross Entropy Loss

$$H(p, q) = - \sum_{x \in \text{classes}} p(x) \log q(x)$$

True probability distribution  
(one-shot)

Your model's predicted  
probability distribution

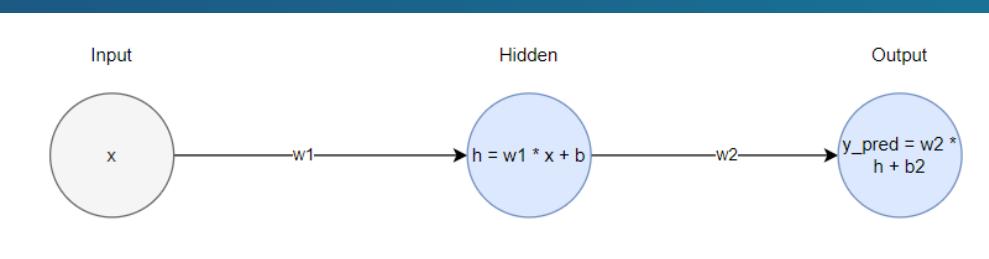
Image source: <https://www.v7labs.com/blog/cross-entropy-loss-guide>

Steps for training a ANN with supervised learning

- Randomly initialize  $\mathbf{w}$  and  $\mathbf{b}$ .
- Do a *forward pass* to obtain  $\hat{\mathbf{y}}$ .
- Calculate loss:  $L(\hat{\mathbf{y}}, \mathbf{y})$ .
- Do backpropagation.
- Use the optimizer (e.g.) SGD to improve weights.
- Repeat from 2. until satisfied.

# Neural Networks - Backpropagation

- Is an algorithm to compute the gradient on each network node with respect to the weights
- (Using MSE as loss function)



Steps for training a ANN with supervised learning

1. Randomly initialize  $w$  and  $b$ .
2. Do a *forward pass* to obtain  $\hat{y}$ .
3. Calculate loss:  $L(\hat{y}, y)$
4. Do backpropagation.
5. Use the optimizer (e.g.) SGD to improve weights.
6. Repeat from 2. until satisfied

## 1. Output Layer Gradients:

- Gradient of Loss w.r.t  $w_2$ :  $\frac{\partial \text{Loss}}{\partial w_2} = (y_{\text{pred}} - y_{\text{true}}) \cdot h$
- Gradient of Loss w.r.t  $b_2$ :  $\frac{\partial \text{Loss}}{\partial b_2} = (y_{\text{pred}} - y_{\text{true}})$

## 2. Hidden Layer Gradients:

- Gradient of Loss w.r.t  $w_1$ :  $\frac{\partial \text{Loss}}{\partial w_1} = (y_{\text{pred}} - y_{\text{true}}) \cdot w_2 \cdot x$
- Gradient of Loss w.r.t  $b_1$ :  $\frac{\partial \text{Loss}}{\partial b_1} = (y_{\text{pred}} - y_{\text{true}}) \cdot w_2$

# Neural Networks - Optimizers

- Is the algorithm to update the weights given the gradient from backpropagation
- Cost = Average Loss for entire dataset

Formula:

$$w_{\text{new}} = w - \text{learning\_rate} * \text{gradient}$$

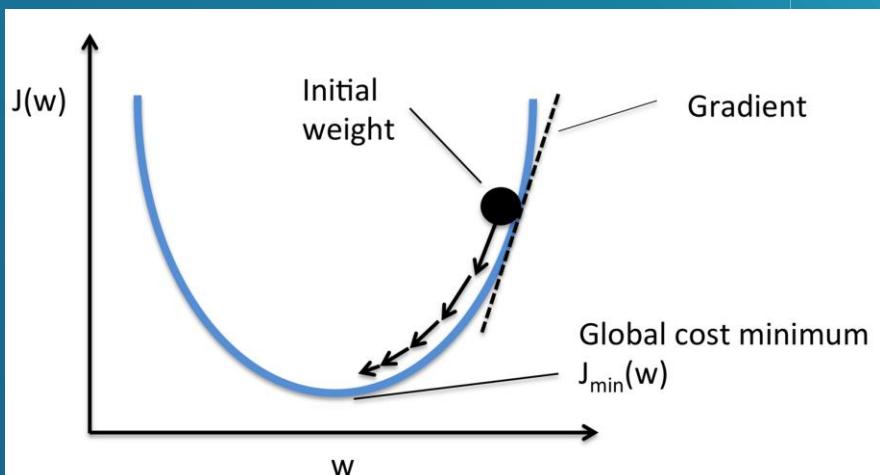
(learning\_rate is a pre-decided hyperparameter)

Image source: <https://www.analyticsvidhya.com/blog/2020/10/how-does-the-gradient-descent-algorithm-work-in-machine-learning/>

Steps for training a ANN with supervised learning

1. Randomly initialize  $w$  and  $b$ .
2. Do a *forward pass* to obtain  $\hat{y}$ .
3. Calculate loss:  $L(\hat{y}, y)$
4. Do backpropagation.
5. Use the optimizer (e.g.) SGD to improve weights.
6. Repeat from 2. until satisfied

## Gradient Descent



# Neural Networks - Optimizers (2)

- Usually Stochastic Gradient Descent is used
- Instead of using the average gradient for the entire dataset, uses the gradient from a small subset of the dataset

```
import numpy as np

# Gradient Descent
def gradient_descent(x, y, theta, alpha, num_iters):
    m = len(y)
    for i in range(num_iters):
        h = np.dot(x, theta)
        loss = h - y
        gradient = np.dot(x.T, loss) / m
        theta = theta - alpha * gradient
    return theta

# Stochastic Gradient Descent
def stochastic_gradient_descent(x, y, theta, alpha, num_iters):
    m = len(y)
    for i in range(num_iters):
        for j in range(m):
            h = np.dot(x[j], theta)
            loss = h - y[j]
            gradient = x[j].T * loss
            theta = theta - alpha * gradient
    return theta
```

Code source:  
[https://medium.com/@dhirendrachoudhary\\_96193/what-is-the-difference-between-stochastic-gradient-descent-sgd-and-gradient-descent-gd-10e7d8019018](https://medium.com/@dhirendrachoudhary_96193/what-is-the-difference-between-stochastic-gradient-descent-sgd-and-gradient-descent-gd-10e7d8019018)

Steps for training a ANN with supervised learning

1. Randomly initialize  $w$  and  $b$ .
2. Do a *forward pass* to obtain  $\hat{y}$ .
3. Calculate loss:  $L(\hat{y}, y)$
4. Do backpropagation
5. Use the optimizer (e.g.) SGD to improve weights.
6. Repeat from 2. until satisfied

# Neural Networks - Optimizers (3)

- Other Optimizers
  - Adam
  - AdamW
- Converge faster than SGD by using Adaptive learning rate  
AdamW has weight decay (to discourage large weights)

Steps for training a ANN with supervised learning

1. Randomly initialize  $\mathbf{w}$  and  $\mathbf{b}$ .
2. Do a *forward pass* to obtain  $\hat{\mathbf{y}}$ .
3. Calculate loss:  $L(\hat{\mathbf{y}}, \mathbf{y})$
4. Do backpropagation
5. Use the optimizer (e.g.) SGD to improve weights.
6. Repeat from 2. until satisfied



# CNN training loop in PyTorch

- Let's see some code



5

# Specific CNN Architectures

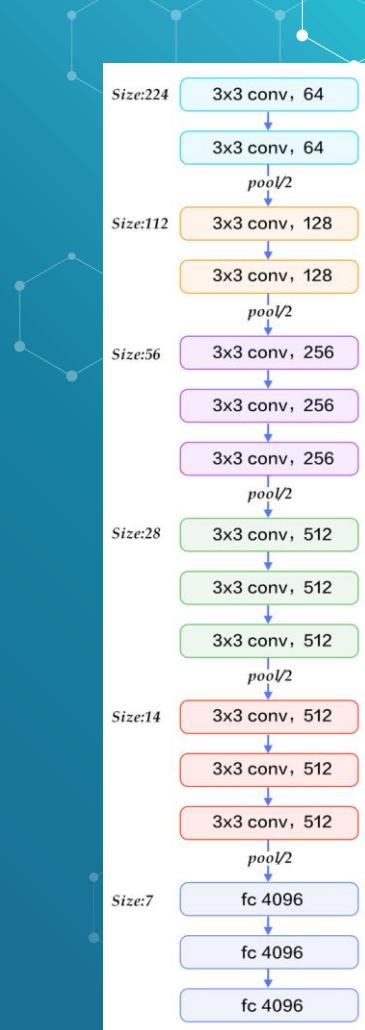
The most successful ones



# VGG

- Inputs  $224 * 224$  pixel RGB images
  - For ImageNet they cropped out the center  $224*224$  patch of the image
- Composed of 16 layers
  - 13 Convolutional  $3 \times 3$  layers
  - Ends with 3 fully connected layers
  - Uses ReLU as the activation function

Image source: <https://www.kaggle.com/code/blurredmachine/vggnet-16-architecture-a-complete-guide>

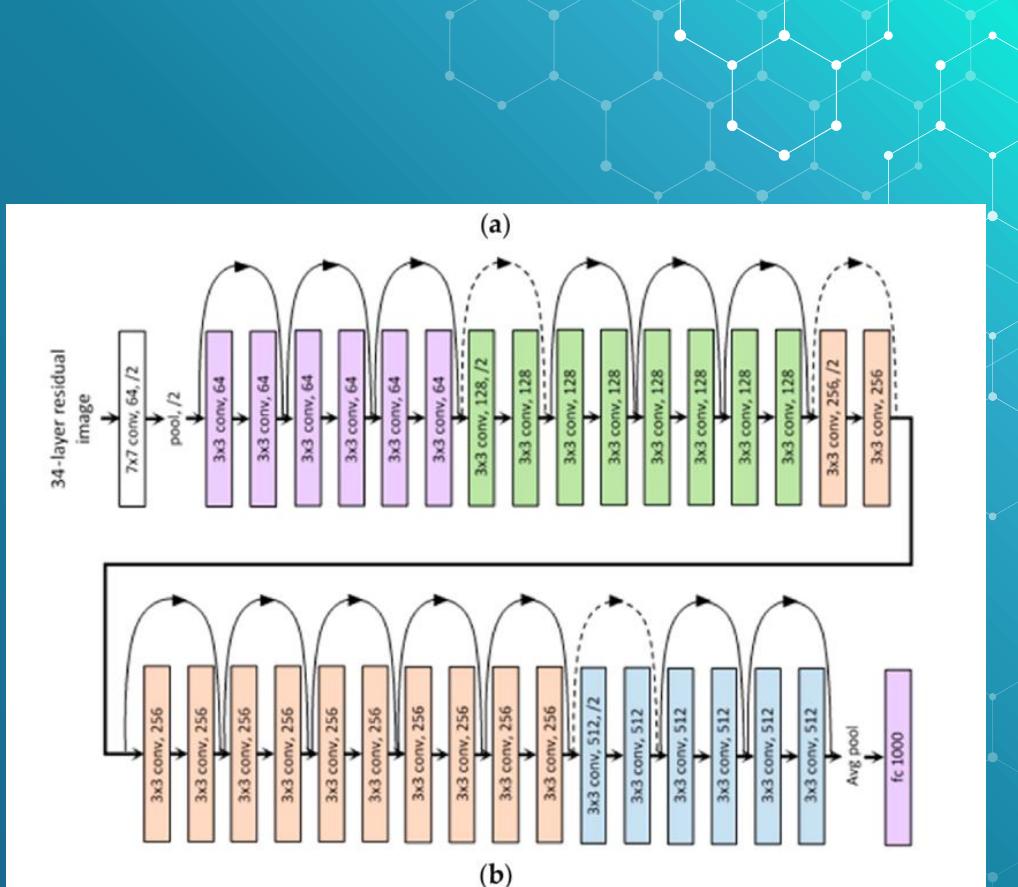




# ResNet

- Introduces the residual connection
  - Skips layers and adds a previous input to the current output.
- Solves the vanishing gradients problem of deep neural networks

Image source: <https://medium.com/@siddheshb008/resnet-architecture-explained-47309ea9283d>  
ResNet paper: <https://arxiv.org/pdf/1512.03385.pdf>

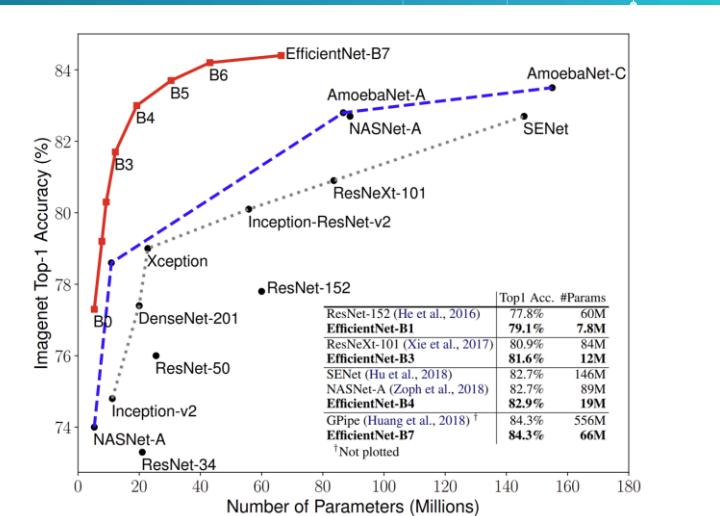


# EfficientNet

- A CNN that builds upon MobileNetV2 and uses MBConv layers
- Introduces an optimal scaling of depth (number of layers) and width (number of neurons per layer) and image resolution by accuracy and desired parameters

The most efficient CNN if you want to use one

Image source and EfficientNet paper: [arxiv.org/pdf/1905.11946.pdf](https://arxiv.org/pdf/1905.11946.pdf)



# CIFAR10

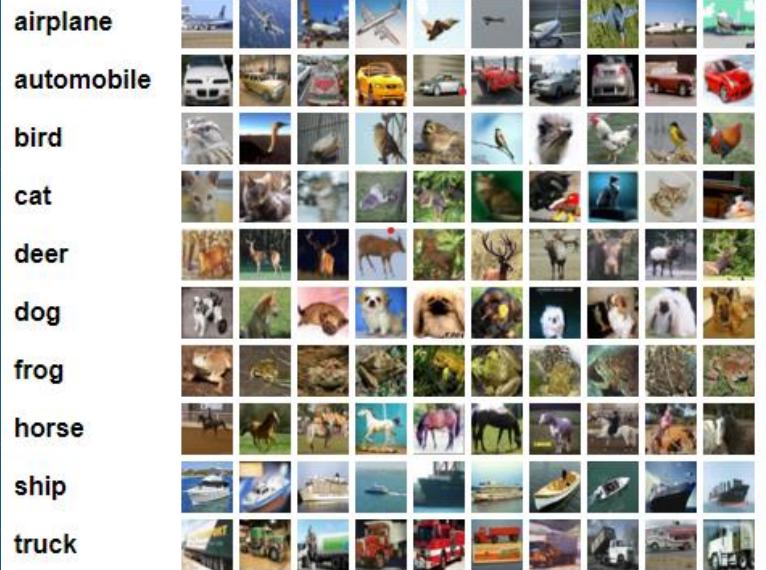


Image source: <https://www.cs.toronto.edu/~kriz/cifar.html>

# COMPETITION!



# Competition: Accuracy on CIFAR10

- **Goal:** Train and validate a custom CNN on the CIFAR10 dataset

## Rules

1. You cannot use pretrained or prebuilt networks, like ResNet50 from torchvision
2. Use the last 20% of the dataset for testing. This data cannot be trained on. (See `cifar\_test` and `cifar\_trainable`)
3. You can only train on the trainable CIFAR10 data, and no other data.
4. You can create as big network as you choose using anything from torch.nn, as long as you build it yourself.
5. You have to complete the challenge within the time limit.